CHIEF RED BEAR CHILDREN'S LODGE ASK RED BEAR DEVELOPMENT STATUS & COMPLETION ASSESSMENT

Evidence-Based Evaluation of Frontline Speech Transcription vs AI Answer Generation
Revision 2 — Clear Operational Differentiation & Sovereign Completion Roadmap

Client Organization: Chief Red Bear Children's Lodge / Cowessess First Nation, Saskatchewan, Canada
Technical Baseline: Post-Phase 5 Hardening | **Git Commit HEAD:** 0746579 | **Branch:** main
Repository Path: C:\Users\Shilley Pc\crbcl | **Remote:** https://github.com/innovatorledger-web/crbcl
Audience: CRBCL Board of Directors, Executive Leadership, IT Steering Committee, and Independent Evaluators

---

■■ ASK RED BEAR ARCHITECTURAL DISTINCTION

Critical Architectural Distinction Notice for CRBCL Leadership: To eliminate any misunderstanding regarding the development status of 'Ask Red Bear', this assessment explicitly separates two fundamentally distinct technical capabilities: (1) Frontline Speech-to-Text Transcription, which converts spoken voice audio into an editable text string inside the user's browser; and (2) AI Answer Generation, which parses natural language questions, searches an organizational knowledge base, and synthesizes intelligent answers. The former is LIVE VERIFIED on the production website; the latter is STUBBED / NOT LIVE via a simulation provider (FakeAiProvider) and has never been connected to a live cloud large language model (LLM).

1. EXECUTIVE SUMMARY & DEVELOPMENT REALITY

Previous documentation created the impression that Ask Red Bear was an operational artificial intelligence assistant capable of autonomous child welfare practice consultation. A rigorous source-code audit reveals the precise technical reality:

• **Speech-to-Text Dictation is fully implemented and live-verified.** When a caseworker clicks the microphone button, audio is recorded, streamed to the backend transcription endpoint, converted into text, and placed into the textarea. No audio recordings or text transcripts are permanently saved to disk, preserving client confidentiality.

• **Artificial Intelligence Answer Generation is stubbed in source code.** In backend/app/services/ai_gateway.py, the active provider is configured as FakeAiProvider. When a user submits a question, the backend returns simulated pre-canned text. There is no active API connection to OpenAI, Anthropic, or any other cloud LLM on production.

• **Operational Impact:** Caseworkers currently possess a highly useful voice dictation tool for drafting progress notes and inquiries, but they do NOT possess an intelligent policy advisory assistant. Full implementation requires Board approval of cloud AI terms, knowledge base curation, and RAG vector infrastructure.

2. DUAL-SUBSYSTEM ARCHITECTURE & DATA FLOW

The diagram below illustrates the exact separation between the live speech transcription subsystem and the stubbed AI generation gateway:

3. COMPARATIVE ANALYSIS: SPEECH DICTATION VS AI GENERATION

The table below compares the two subsystems across all operational, technical, and governance criteria:

Evaluation Dimension	Speech Dictation (Subsystem A)	AI Answer Generation (Subsystem B)
Core Primary Purpose	Frontline dictation aid to accelerate typing case notes and intake summaries.	Intelligent policy advisory assistant answering practice and governance questions.
Current Physical Status	LIVE VERIFIED on live website (https://genserver.online).	STUBBED / NOT LIVE in backend (routes to FakeAiProvider).
Data Persistence Standard	Zero persistence; ephemeral in-memory audio buffer discarded immediately.	Proposed immutable audit log of policy inquiries and retrieved citations.
External Cloud Dependency	Whisper STT service (local container or fast API bridge).	Requires sovereign enterprise API key for Board-approved cloud LLM provider.
First Nations Privacy Risk	Very low; audio is never stored; text is explicitly reviewed by caseworker.	Moderate to high; requires strict prompt boundaries and zero-data-retention agreements.
Caseworker Autonomy	High; caseworker maintains 100% editorial control over typed words.	Advisory only; answers must cite specific Cowessess Act sections.
Readiness for Daily Work	100% operationally ready and currently accessible to staff.	Requires 4-6 weeks of knowledge curation, vector indexing, and supervisor testing.

4. DETAILED AUDIT OF 12 ASK RED BEAR CAPABILITIES

Below is an audited inventory of all 12 frontend, backend, security, and infrastructure capabilities associated with the Ask Red Bear module:

Capability Name	Implementation Details	Current Status	Operational Scope & Evidence
Frontend Chat Widget & Modal	AskRedBearModal.jsx, FloatingRedBear.jsx	**LIVE VERIFIED**	Responsive overlay modal, expandable drawer, conversation history UI, and input state transitions.
Microphone Web Audio Capture	SpeechRecorder.jsx, audio_utils.js	**LIVE VERIFIED**	Browser MediaStream capture, 16kHz audio downsampling, WebM/WAV encoding, visual VU meter waveform.
Speech-to-Text Transcription	POST /api/v1/ask-red-bear/transcribe	**LIVE VERIFIED**	FastAPI transcription bridge; ephemeral audio processing; tested operational on https://genserver.online.
Transcribed Text In-Place Editing	Textarea event binding in React	**LIVE VERIFIED**	Spoken words stream into editable textarea; caseworker can manually correct names and Cree terminology.
Audio Privacy & Ephemeral Buffering	In-memory BytesIO streaming	**LIVE VERIFIED**	Zero audio disk persistence; audio stream discarded immediately following text transcription extraction.
AI Gateway Abstraction Interface	backend/app/services/ai_gateway.py	**SOURCE IMPLEMENTED**	Clean provider interface pattern (AiProvider base class) supporting interchangeable LLM backends.
FakeAiProvider Simulation Engine	app/services/ai_gateway.py:FakeAiProvider	**STUBBED / NOT LIVE**	Deterministic mock response generator; returns simulated answers for local integration testing.
Live Cloud LLM Inference Engine	Candidate Cloud LLMs (Claude, Azure OpenAI, Sovereign Model)	**PLANNED / NOT IMPLEMENTED**	Direct integration with commercial or sovereign cloud LLM API; pending Board evaluation, selection, and API key procurement.
RAG Vector Document Indexing	PostgreSQL pgvector extension	**PLANNED / NOT IMPLEMENTED**	Vectorized knowledge store of Cowessess Miyo-Pimatisiwin Act, CRBCL practice manuals, and policies.
Role-Aware Context Blinding	RAG security filter middleware	**PLANNED / NOT IMPLEMENTED**	Enforcement of RBAC permissions and Case Restrictions to ensure AI never leaks confidential family data.
Prompt Injection & Guardrails	Input sanitization & boundary prompts	**PLANNED / NOT IMPLEMENTED**	System prompt hardening preventing adversarial extraction of system instructions or protected records.
First Nations OCAP® Data Sovereignty	Canadian sovereign data hosting boundary	**DECISION PENDING**	Contractual and technical guarantees that CRBCL queries and embeddings are never used for model training.

5. TECHNICAL ARCHITECTURE OF AI_GATEWAY.PY AND FAKE_AI_PROVIDER

In backend/app/services/ai_gateway.py, the system implements an abstract provider interface pattern:

```
class AiProvider(ABC):
    @abstractmethod
    async def generate_response(self, prompt: str, context: Optional[str] = None) -> AiResponse:
        pass

class FakeAiProvider(AiProvider):
    async def generate_response(self, prompt: str, context: Optional[str] = None) -> AiResponse:
        # Returns simulated response for testing
        return AiResponse(
            text=f"This is a simulated Ask Red Bear response to: '{prompt}'. Real AI integration is pending.",
            tokens_used=12,
            model="simulation-stub-v1"
        )
```

Because FakeAiProvider is hardcoded as the active instance, all POST requests to /api/v1/ask-red-bear/query return static mocked responses. To activate real AI, an AnthropicClaudeProvider, AzureOpenAiProvider, or Canadian sovereign model class must be instantiated and injected upon Board approval.

6. PROPOSED RETRIEVAL-AUGMENTED GENERATION (RAG) COMPLETION PIPELINE

To transform Ask Red Bear from a voice dictation tool into an authoritative sovereign policy assistant, the following 6-step engineering pipeline must be executed:

• **1. Approved Knowledge Ingestion:** CRBCL policy directors provide authoritative PDF/DOCX documents: Cowessess Child & Family Services Act, Miyo-Pimatisiwin Act regulations, practice guidelines, supervisor manuals, and standard operating procedures. Documents are cryptographically hashed and versioned.

• **2. Document Chunking & Vectorization:** Text is segmented into 500-token semantic chunks with 100-token overlap. Each chunk is passed to a text-embedding model (e.g. text-embedding-3-small or self-hosted BGE-large) generating 1536-dimensional dense vectors.

• **3. Sovereign Relational Storage:** Embeddings are stored in CRBCL's managed PostgreSQL database using the pgvector extension. Each chunk is tagged with metadata: document_id, section_name, effective_date, required_role, and classification_level.

• **4. Semantic Retrieval & RBAC Filtering:** When a user submits a question, their query is vectorized. An approximate nearest neighbor (ANN) cosine similarity search retrieves the top-k most relevant policy chunks. A strict RBAC filter purges any chunks exceeding the user's role authorization.

• **5. Context Construction & System Prompting:** The retrieved chunks are injected into a hardened system prompt: 'You are Ask Red Bear, an internal knowledge assistant for Chief Red Bear Children's Lodge. Answer the caseworker's question strictly using the provided Cowessess First Nation policy context below. If the context does not contain the answer, state that you do not know and advise consulting a supervisor. Never invent policies.'

• **6. Cloud Inference & Audit Logging:** The prompt is dispatched via TLS 1.3 to an approved cloud LLM provider (configured with zero-data-retention terms). The generated response is returned to the user, and an immutable record (user_id, timestamp, token_count, retrieved_chunk_ids) is logged for quality audits.

7. FIRST NATIONS OCAP® PRINCIPLES & DATA SOVEREIGNTY SAFEGUARDS

Integrating artificial intelligence into a First Nations child welfare agency requires strict adherence to OCAP® (Ownership, Control, Access, and Possession) principles:

• **Ownership:** Cowessess First Nation retains exclusive intellectual property and ownership over all ingested policies, sacred teachings, and operational data.

• **Control:** CRBCL leadership controls which documents are ingested, which roles can query them, and what guardrails govern AI responses.

• **Access:** Frontline staff access is gated by the 327-permission RBAC engine. Caseworkers cannot access information regarding cases where active Case Restrictions exist.

• **Possession:** Data must reside on Canadian sovereign infrastructure. Enterprise zero-data-retention agreements must ensure commercial AI vendors never train on CRBCL data.

8. TRANSPARENT CLOUD LLM OPERATING COST PROJECTIONS

To prevent budget surprises, operational token costs have been modeled based on typical enterprise tier rates (e.g., Anthropic Claude 3.5 Sonnet benchmark pricing of $3.00 / million input tokens, $15.00 / million output tokens; final costs depend on Board vendor selection):

• **Frontline Staff Assumption:** 40 active caseworkers and supervisors conducting an average of 5 policy inquiries per business day = 200 inquiries/day (4,200 inquiries/month).

• **Token Consumption per Query:** Average 1,500 prompt tokens (user question + retrieved policy context chunks) + 400 completion tokens (synthesized answer).

• **Monthly Token Volume:** 6.3 million input tokens ($18.90 CAD) + 1.68 million output tokens ($25.20 CAD).

• **Vector Index Hosting (PostgreSQL pgvector):** $0.00 additional cost (runs natively inside CRBCL's existing managed database container).

• **Total Projected AI Operating Cost:** Approximately **$45.00 to $75.00 CAD per month** for the entire organization under regular operational utilization.

9. PRIORITIZED COMPLETION ROADMAP & GOVERNANCE DECISIONS

To operationalize AI answer generation, the following governance milestones must be achieved:

1. **Board Decision:** Formal approval of enterprise cloud AI terms of service (zero-data-retention agreement).
2. **Knowledge Curation:** Executive Director sign-off on the official corpus of policy PDFs for vector ingestion.
3. **Infrastructure Deployment:** Enablement of pgvector extension on Railway PostgreSQL database.
4. **Security Review:** Execution of automated prompt injection and case conflict blinding test suites.
5. **Supervisor UAT:** Structured evaluation workshops with frontline supervisors to test answer accuracy.